

Data Synchronization and Replication Tool

Pradeeban Kathiravelu

INESC-ID Lisboa

Instituto Superior Técnico, Universidade de Lisboa

Lisbon, Portugal

Email: pradeeban.kathiravelu@tecnico.ulisboa.pt

Ashish Sharma

Department of Biomedical Informatics

Emory University

Atlanta, Georgia, USA

Email: ashish.sharma@emory.edu

Abstract—Consumers download the data by searching the image repository using the browser. The information that the consumer is interested in, gets updated whenever the data producers update or add patient information. The current download tool lacks the ability to track the relevant updates to the consumer. A data replication and synchronization tool will assist automated downloads to the consumers. Users can create replica sets as a sub set of their search queries, and share their replica sets with other users, and update their replica sets periodically. Replica sets can be used as a way of tracking and sharing information. This project exploits this model to create a data replication and synchronization tool for data sources. As a proof of concept, a data replication and synchronization tool has been implemented for The Cancer Imaging Archive (TCIA).

I. INTRODUCTION

Data sources contain data of different granularity. Data is organized in a hierarchical structure, where different levels are used to present data in specific formats. Folders and documents make a good example of this. A unique identifier is assigned to each of the data units. A search across the data source would present the user with the list of matching criteria. Interesting sub set of the matching criteria may be bookmarked by the user and shared with others.

The sub set which is commonly known as a replica set, can be updated, duplicated, shared with other users, and deleted later. A search query may contain different parameters that can define the scope of the search and the outcomes, and often return the outputs in a finer granularity. When a sub set of such information is stored as a replica set, it is sufficient to store the unique identifiers of the matching data units of finer granularity than the original search query, as it would be sufficient to reproduce the data that is represented by the replica set.

II. BACKGROUND

A. Representation of Medical Images in TCIA

Medical images are represented in multiple granularity. Figure 1 represents how the images are structured hierarchically in TCIA.

III. DESIGN AND IMPLEMENTATION

Having multiple instances running over different nodes provide fault-tolerance, as when one node terminates, the other nodes have the backup replica of the partitions stored in the terminated node. Figure 2 shows the higher level deployment view of the solution.

A. Generic Design

Two distributed cache instances exist in InfDataAccessIntegration.

```
protected static Cache userReplicasMap;
protected static Cache replicaSetsMap;
```

userReplicasMap is a mapping of userID → Array of replicaSetIDs. UserID could be the logged in user name. (for now, testing with random strings). replicaSetsMap is a mapping of replicaSetID → replicaSet.

Though this could be replaced with a single cache instance with the mapping of userID → replicaSets, having two cache instances will be more efficient during searches, duplicates, and push changes. Hence, two cache instances design was chosen.

InfDataAccessIntegration implements the PubConsAPI for publisher/consumer, where *InterfaceManager* implements the *InterfaceAPI* for the interface between the data source and the data replication and synchronization tool. Invoker classes extending the abstract class *InterfaceManager*, implement the respective data source integration to invoke these




Fig. 1. Medical Images Granularity




Fig. 2. Deployment

methods. The execution flow is depicted by Figure 3. When the user logs in, *login()* checks whether the user has already stored *replicaSets* from the Infinispan distributed Cache. If so, execute them all again. This would be changed later as we do not have to execute all. Rather, we need to execute for the diffs. When the user performs new searches, for the images, series, collections, and the other meta data, the results will be returned to the user, and the user can chose ta subset of the returned results to create a *replicaSet*. The *replicaSet* for the

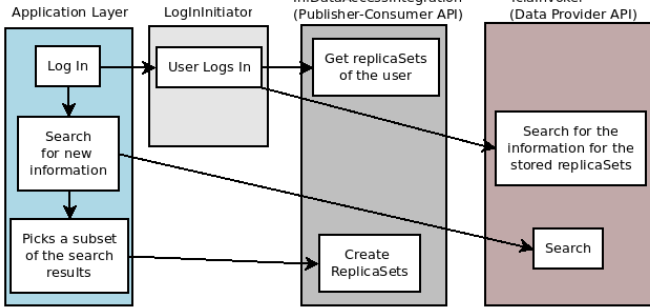


Fig. 3. Execution Flow

image will be as,

```
TCIAConstants.IMAGE_TAG + "getImage?SeriesInstanceUID=" +
    seriesInstanceUID
```

For other information (meta data), such as collections and seies,

```
TCIAConstants.META_TAG + query;
```

Here, query takes the below format.

```
"getSeries?format=" + format +
    "&Collection=" + collection +
    "&PatientID=" + patientID +
    "&StudyInstanceUID=" + studyInstanceUID +
    "&Modality=" + modality;
```

When a new instance starts now, and invokes the log in action for the same user, it will execute the queries for the stored *replicaSets* again, and reproduce the same results.

B. Design and Implementation for TCIA

Different complex data sources require custom development extending the generic framework. As creating and customizing the *replicaSet* require a more specific data structure, further implementations are done, extending the core class hierarchy. An extension based on the base design was developed for TCIA, as shown by Figure 4, which provides a core class hierarchy of the system.

TciaInvoker - Interface API: *TciaInvoker* extends *InterfaceManager* to implement the interfacing layer between the TCIA data source and the data replication and synchronization tool. Meta data such as collections, patients, studies, and series are retrieved at different levels, though the default download manager of TCIA downloads the data in series level, composed of the images of the series in a single zip archive. While having the default *userReplicasMap* to contain

the IDs of the replica sets for each user, the replica set itself is stored in multiple maps instead of a single *replicaSets* map, to provide an efficient storage and access.

DataProSpecs extends the *InfDataAccessIntegration* class. 5 maps are created as below to represent the replica sets.

```
protected static Cache<Long, Boolean[]> tciaMetaMap;
protected static Cache<Long, String[]> collectionsMap;
protected static Cache<Long, String[]> patientsMap;
protected static Cache<Long, String[]> studiesMap;
protected static Cache<Long, String[]> seriesMap;
```

tciaMetaMap contains a boolean array, which reflects which of the granularity of meta data is selected as a whole. For TCIA, if a few collections are selected, the first element of the array is set to true, and similarly, the other meta data are marked to true or false as shown by the below code segment.

```
Boolean[] metaMap = new Boolean[4];
metaMap[0] = collection != null;
metaMap[1] = patientID != null;
metaMap[2] = studyInstanceUID != null;
metaMap[3] = seriesInstanceUID != null;

putReplicaSet(replicaSetId, metaMap);
```

The name of the collections, *patientID*, *studyInstanceUID*, and *seriesInstanceUID* are stored against the respective *replicaSetID* in *collectionsMap*, *patientsMap*, *studiesMap*, and *seriesMap* respectively. Hence changes are done at the respective maps. Duplicating the *replicaSets* duplicate the contents of the entire row to a new *replicaSetID*. Similarly, deleting a *replicaSet* deletes the respective information from all the maps.

TCIA public API provides methods to retrieve the images and meta data of different granularity. Figure 5 depict the methods that retrieve image and metadata at different granularity from TCIA. These methods are invoked by the replication manager tool to retrieve the images. As shown by Figure 5 an initial search on TCIA may contain parameters such as modality, in addition to collection name, patient ID, study instance ID, and series instance UID. However, each of these search returns the output in a finer granularity, and when a sub set of this finer granularity is selected, each of the selected elements are always identified by their respective identifier. Hence, storing an array of patient ID would be sufficient to identify the selected sub set of the collection including the array of patients. Similary an array of study instance UID is sufficient to represent the selected sub set of any patient, and an array of series instance UID is sufficient to represent the selected sub set of any study, as it will contain an array of series.

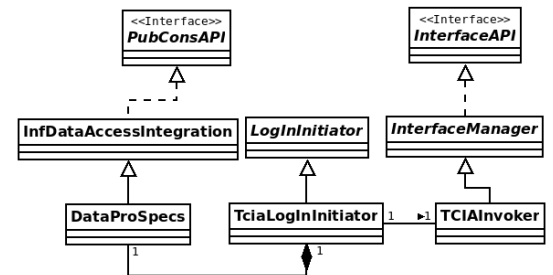


Fig. 4. Core Class Hierarchy




Fig. 5. Retrieving images and meta data

multi-tenancy: The Replication Tool is multi-tenanted, and it is aware of the multiple tenants or users using the system. Each user co-exist in the replication tool at the same time, without the knowledge of existence of the other users, sharing the same cache space.

DataProSpecs - PubConsAPI: The methods of *DataProSpecs* can be invoked by knowing the respective *replicaSetID* of the replica set and the user ID of the user who owns the replica set. User ID is a random string that is input by the user. Probably this can be some 'secret' or pass from the user. userIDs serve as the keys of the *userReplicasMap*. Apart from that, there is no data structure to hold the list of users.

createReplicaSet() requires the *userID* along with the elements to be stored in the replica set. It returns the *replicaSetID*, which is further used to uniquely identify the replica set. *getReplicaSet()* requires only the relevant *replicaSetID* to return the respective replica set. Similarly, *updateReplicaSet()* requires *replicaSetID* as well as the elements to be stored in the *replicaSet*, replacing the previous elements. *deleteReplicaSet()* requires both the *userID* and *replicaSetID*. Similarly, *duplicateReplicaSet()* requires both *replicaSetID* and the *userID* of the user to whom the replicaSet is shared to. Creating, deleting, and duplicating a replication set requires modification to the user replicas map, which holds the list of the replica sets for the particular user. Hence the necessity to input the *userID*. *ReplicaSetID* is a large negative number (UUID) randomly generated. It can be safely assumed to be harder to guess. Hence, this model is assumed to provide adequate security for this prototype.

C. Integration with Multiple Data Sources

Clinical data is deployed in multiple data sources such as TCIA, CA Microscope, and Amazon S3. Figure 6 depicts the deployment of the system with multiple data sources. A set of clinical data was uploaded to S3, where the meta data mapping of *patientID* → *respectivefilename* was available as a CSV file. Similarly CSV file depicting clinical information is available, as multiple properties against the UUID, such as patient ID. These files are parsed and stored into meta data maps in Infinispan. The CSV files containing meta data or *resourceID* → file name mapping are stored locally in the file system or in a remote repository.

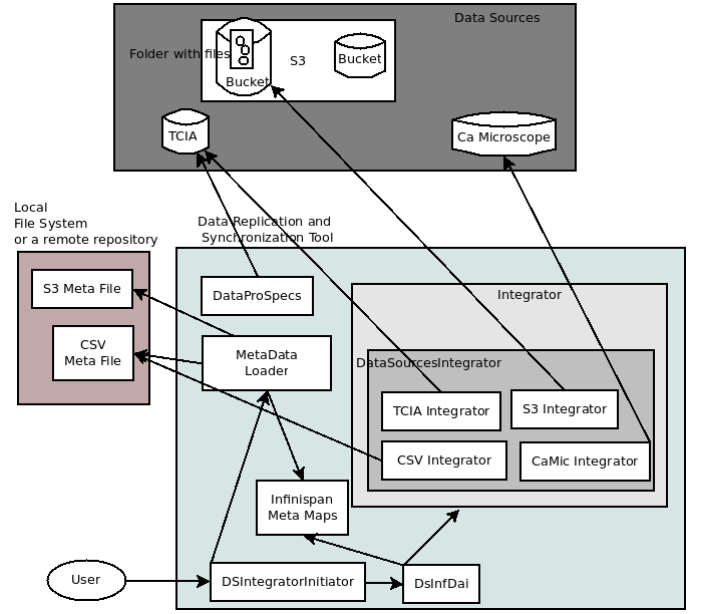


Fig. 6. Deployment Diagram of the System and the Data Sources

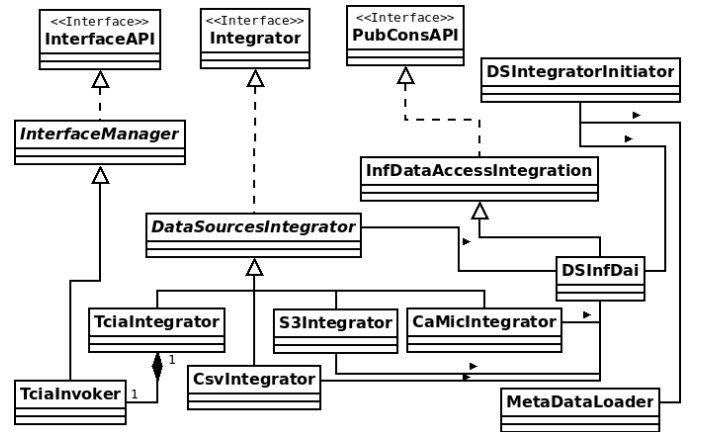


Fig. 7. Class Diagram of data sources integrators along with the higher level API

IV. EVALUATION

A prototype web application was built with the data replication and synchronization tool and TCIA REST API. Apache Velocity was used to generate the web pages for the replication tool. Apache Tomcat (Embedded) was integrated into the program such that it will get the user inputs from the HTML pages and present the output in pages formatted by Apache Velocity Templates. Respective servlets were created inside the servlets package to receive the inputs from HTML pages to the backend Java code.

DataRetriever class is invoked to retrieve the meta data from TCIA and to create and retrieve replica sets. *UIGenerator* class invokes the Apache Velocity templates to create the web pages, presenting the data to

the users. *DataRetriever* has instances of *TciaInvoker* and *TciaLogInInitiator* to invoke the TCIA queries and replication tool create and retrieve replicaSets mechanisms. *TciaLogInInitiator* has instances of *TciaInvoker* and *DataProSpecs*, to invoke the public APIs provided by these classes for the data provider specification and publisher/consumer APIs. Figure 7 depicts the class hierarchy involved in the user interface.




Fig. 8. Class Hierarchy of the User Interface

V. CONCLUSION

Storing a selected sub set of data items that matches a specific criteria as a replica set is an optimal way to bookmark large data sources. Users should be able to create and update their replica sets, and share it with others using a unique identifier. While different data sources have different interfaces, a generic data replication and synchronization tool will be convenient, such that sharing of replica sets across heterogeneous data sources will be possible. This project researched the possibilities of such a replication and data synchronization tool, exploiting the in-memory data grid projects, and implements a data synchronization and replication tool¹ with Infinispan while consuming data sources such as TCIA via their public APIs.

VI. FUTURE WORK

Get the Raw Images: `downloadRawData(Stringkey)` method should be implemented in the classes implementing the

interface to be able to download the raw data from the data sources such as TCIA, CA Microscope, or S3, for the given replica set, potentially using Java web start.

Tracking the downloaded items: Currently, when a user initiates download of a replica set, finishes the download, and later restart the download, the download mechanism still would download the whole replica set again, as there is no track of what has already been downloaded. A sample timeline of the events is given below.

t_0 : User A Creates Replica Set R_1 with a query.

t_1 : User A accesses and downloads the contents of the replica set R_1 and finds the contents to be A, B, C .

t_2 : The contents matching the replica set R_1 query modified to be A, B, D by an outside event such as more images added and existing images removed.

t_3 : User A downloads the contents of R_1 again.

This should be designed such that the already accessed images would not be downloaded while the newly available or not previously accessed segments of the replica set will be downloaded. For example, in the above scenario, currently at

¹The source code can be accessed from <https://bitbucket.org/BMI/datareplicationsystem>

t_3 , all A, B, and D will be downloaded, where an implementation should allow only the new or modified contents to be downloaded.

Security: The data replication and synchronization can be further secured with user authentication. SAML tokens can be created for user names and passwords with the membership information. The SAML tokens can be validated later. A create and validate API with user stores such as LDAP or OpenID should be designed and implemented, extending and leveraging the replication tool.

A. Developer Guide

This section describes how the future work could be done.